





New Directions in Hypervisor Detection

by cristianrichie



About Me

Cristian Assaiante [[@cristianrichi3](https://twitter.com/cristianrichi3)]

- 23 years old, student of Engineering in Computer Science.
- Reverse engineering, Binary Exploitation and Compilers.
- Capturing flags with TRX and mhackeroni (since 2018).



(stock image from DEFCON27)



!Disclaimer!

- **My Ticks Don't Lie:** New Timing Attacks for Hypervisor Detection
 - BlackHat EU 2020
 - Daniele Cono D'Elia
 - [\[original talk\]](#)

My Ticks Don't Lie: New Timing Attacks for Hypervisor Detection

[Daniele Cono D'Elia](#) | Postdoctoral researcher, Sapienza University of Rome

Date: Thursday, December 10 | 11:20am-12:00pm

Format: 40-Minute Briefings

Tracks:  Malware,  Cloud & Platform Security

Hypervisor detection is a pillar of sandbox evasion techniques. While hardware-assisted virtualization solutions are indispensable for scalable dynamic malware analysis, compared to bare-metal machines they all introduce timing discrepancies that expert malware writers may reveal using low-level measurement sequences. Today, the most advanced sandboxes fight such attempts by massaging the values malware can read from classic time probes. We will see how this battle is far from over: by taking advantage of recent developments in hardware-assisted virtualization, we will build and exercise two novel hypervisor detection primitives that current sandboxes are incapable of handling. The first idea is to build a high-resolution TSC counter, of which we will show a thread that can tick just as accurately as an unpatched TSC counter, of which we will show a well-known attack on the last-level cache to detect pollution caused by the hypervisor. We will also show that while several classic time evasions have been patched back to life without raising alerts, new ones have appeared on their radars, and may



➔ **Hypervisor Detection 101**

Covert Time Source

Retrofitting Red Pills

LLC Prime+Probe Attack



Malware Analysis and Virtualization


- Hypervisors cannot be avoided
 - analysts use VMs to safely perform their tasks
 - sandboxes run VMs on servers
- Hypervisor detection is a pillar of malware evasion techniques



Virtualization 101

VMX (Virtual Machine eXtension) enables CPU support for virtualization

- **VMM** (Virtual Machine Monitor) acts as host: retains selective control of hardware resources and offers virtual processors to guests
- VMM runs in VMX root mode, while guest in non-root mode




Interaction of a VMM and Guests
(from: Intel 64 and IA-32 Architectures SDM)



Virtualization 101

During the execution in guest environment, upon performing certain tasks, a transition to VMM occurs

- **VM Exit:** guest state is saved and VMM does its stuff
- **VM Entry:** guest state is restored and guest can continue its execution



Interaction of a VMM and Guests
(from: Intel 64 and IA-32 Architectures SDM)



The CpuID Case

- cpuid (CPU IDentification) instruction allows software to discover information about the running processor
- Input: the cpuid “leaf”, a 32 bit integer read from eax
 - for some leaves, also ecx can be set and it will be treated as input value
- Output: CPU information available within registers eax, ebx, ecx, edx

```
int eax = <leaf>, ebx, ecx, edx;  
__asm__ volatile ("cpuid" :  
                  "a=" (eax), "b=" (ebx), "c=" (ecx), "d=" (edx) :  
                  "a" (eax));
```



The CPUID Case

- cpuid instruction causes a VM Exit event
- Upon it VMM can control exposed properties of the virtual CPU

```
int ecx;  
__asm__ volatile ("cpuid" : "c=" (ecx) : "a"(1) : ...);  
printf("%d\n", (ecx >> 31));
```

31st bit of Extended Feature Information is the hypervisor bit



Timing VM Exit Events

Hypervisor detection via time
comparison of `cuid` execution.

Native: ~300 cycles

VirtualBox 5.2: ~3000 cycles

```
movl $1, %eax  
mfence  
rdtsc  
movl %eax, %esi  
cuid  
rdtsc  
subl %eax, %esi  
negl %esi
```



Sandboxes Anti-Evasion Tricks

- Track instructions causing VM Exit
- Optimize VMM code to reduce VM Exit time overhead
- Fake time sources return values:
 - rewrite output of some APIs
 - make `rdtsc` cause a VM Exit too, then alter its returned values keeping track of time spent in VMM

How can we measure time without any available time source?



Hypervisor Detection 101

➔ **Covert Time Source**

Retrofitting Red Pills

LLC Prime+Probe Attack



Building a Covert Time Source

```
volatile uint64_t ctr_clock;

// spawn a cthread
while (1) {
    ctr_clock++;
}

// main code
uint64_t start, end;
start = ctr_clock;
__asm__ volatile ("cpuid" ...);
end = ctr_clock;
```

Issues:

- fast enough?
- reliable?
- serialization?

Approximate Resolution:

- rdtsc: ~2793.5 Mhz
- cthread: ~540 MHz
 - w/ TurboBoost



A Clever Implementation

```
volatile uint64_t ctr_clock;
```

```
// spawn a cthread
```

```
__asm__ volatile(
```

```
    "xorq %%rax, %%rax;"
    "movq %0, %%rcx;"
    "1: incq %%rax;"
    "    movq %%rax, (%%rcx);"
    "jmp 1b"
```

```
:
```

```
: "r"(&ctr_clock)
```

```
: "rax", "rcx"
```

```
);
```

Trick: avoid reading counter value from memory to update it

Why: cost of L1 access time impacts update frequency; inc and mov have good latency and throughput

From: "Malware Guard Extension: abusing Intel SGX to conceal cache attacks" by Schwarz, Weiser, Gruss, Maurice, Mangard. Springer Cybersecurity 2020.



A Clever Implementation

```
volatile uint64_t ctr_clock;
```

```
// spawn a cthread
```

```
__asm__ volatile(
```

```
    "xorq %%rax, %%rax;"
    "movq %0, %%rcx;"
    "1: incq %%rax;"
    "    movq %%rax, (%%rcx);"
    "jmp 1b"
```

```
:
```

```
: "r"(&ctr_clock)
```

```
: "rax", "rcx"
```

```
);
```

Approximate Resolution:

- rdtsc: ~2793.5 Mhz
- cthread: ~540 MHz
 - w/ TurboBoost
- cthread+: ~**3500** MHz
 - w/ TurboBoost

From: "Malware Guard Extension: abusing Intel SGX to conceal cache attacks" by Schwarz, Weiser, Gruss, Maurice, Mangard. Springer Cybersecurity 2020.



Cores Availability

```
unsigned cnt = 0, i = 0; uintptr_t lst = 0;
for (i = 0; i < LOOP_COUNT; i++) {
    unsigned cur = *(scz->other);
    if (cur == lst) cnt++; else lst = cur;
    __asm__ volatile ( // busy loop
        "xorl %%eax, %%eax;"
        "movl $10000, %%eax;"
        "1: decl %%eax;"
        "jnz 1b"
        : : "eax");
    (*(scz->self))++;
}
```

At least two cores are needed to be able to use counter threads.

How can we check how many cores are available?



Cores Availability

```
unsigned cnt = 0, i = 0; uintptr_t lst = 0;
for (i = 0; i < LOOP_COUNT; i++) {
    unsigned cur = *(scz->other);
    if (cur == lst) cnt++; else lst = cur;
    __asm__ volatile ( // busy loop
        "xorl %%eax, %%eax;"
        "movl $10000, %%eax;"
        "1: decl %%eax;"
        "jnz 1b"
        : : "eax");
    (*(scz->self))++;
}

typedef struct {
    volatile uintptr_t* self;
    volatile uintptr_t* other;
} scz_t;
```



Cores Availability

```
unsigned cnt = 0, i = 0; uintptr_t lst = 0;
for (i = 0; i < LOOP_COUNT; i++) {
    unsigned cur = *(scz->other);
    if (cur == lst) cnt++; else lst = cur;
    __asm__ volatile ( // busy loop
        "xorl %%eax, %%eax;"
        "movl $10000, %%eax;"
        "1: decl %%eax;"
        "jnz 1b"
        : : "eax");
    (*(scz->self))++;
}
```

Race two thread, check sum of two
count variable $< \text{LOOP_COUNT}/2$
(why: count increases when counter
from other thread was not updated)



Hypervisor Detection 101

Covert Time Source

➔ **Retrofitting Red Pills**

LLC Prime+Probe Attack



Detection 1: cpuid latency

Time to execute `cpuid` > threshold

- Set `eax = 1`
- Compute average time
 - $N = 10$ observations
 - threshold = 1000



Detection 2: cpuid/nop ratio

Execution time ratio between `cpuid` and another instruction

- Use a low-latency instruction as reference (e.g. `nop`)
- Idea: different instructions execute similarly slower or faster under different CPUs and conditions
 - Native environment vs Guest environment

Reproduction of “Detecting hardware-assisted virtualization” [DIMVA ‘16]



Detection 3: TLB Eviction

Look for TLB entries evicted by VMM execution

- **TLB** (Translation Lookaside Buffer): keeps track of translations between Virtual and Physical addresses to avoid repeating translations
 - a kind of cache memory for MMU
- Fill TLB and cause a VM Exit. Then analyze memory access latencies

Reproduction of “Detecting hardware-assisted virtualization” [DIMVA ‘16]



Results

76 completed executions on online available sandboxes

- Detection 1: 43/76 detected
 - fake TSC values for 6
 - single core detected for 14
- Detection 2: 55/76 detected
 - fake TSC values for 7
 - high noise for nop for 12 (also with CT)
- Detection 3: 24/29 detected
 - completed only by 52 machines
 - unstable latencies for 23



Hypervisor Detection 101

Covert Time Source

Retrofitting Red Pills

➔ **LLC Prime+Probe Attack**



New Detection

Look for effects on caches quite reliable to measure

Idea:

- VMM execution may evict LLC lines
- Search for those lines to detect VMM

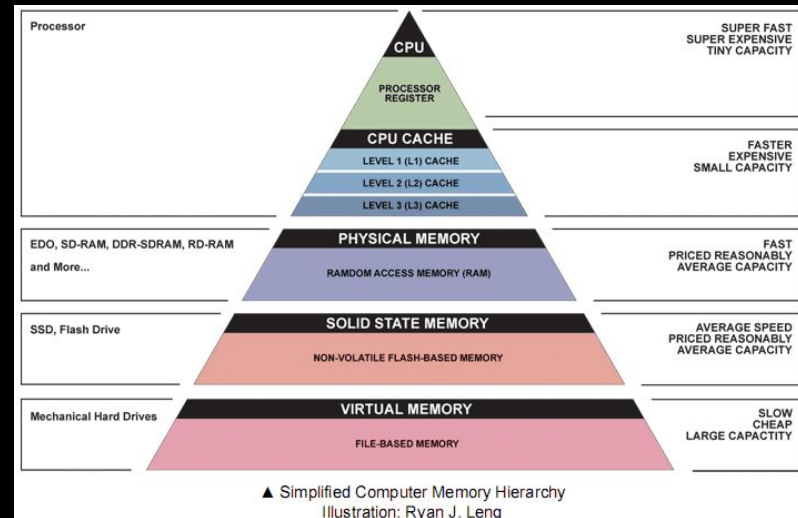
Why LLC?

- High resolution
- Shared between cores




Memory Hierarchy

- Huge difference between CPU and Physical Memory
- To speed up the memory accesses, cache memories are in between
 - L1d/L1i per-core
 - L2 per-core
 - L3 shared
 - a.k.a. **LLC** (last level cache)





LLC Functioning





Prime+Probe Attack



i-th cache set
(16-way associative)



fill each cache set entry (how?)



VMM may evict one or more lines

Some attacker-controlled lines will see higher latency from LLC miss




Eviction Sets

- An eviction set contains virtual addresses that map to one cache set
- Cache associativity determines optimal size
 - $\#ways = \#elements$ of minimal eviction set
- We need to build a minimal eviction set for all available colors




LLC Addressing





Addressing inside the VM



In a sandbox we have no knowledge of the mapping VA-HPA



Finding Minimal Eviction Sets

Theory and Practice of Finding Eviction Sets [S&P19]

- no assumption on the mapping between virtual addresses and cache sets
- $O(n \cdot w^2)$ makes it practical
- does not work with non-inclusive LLC

Theory and Practice of Finding Eviction Sets

Pepe Vila^{1,3}, Boris Köp² and José F. Morales¹

¹IMDEA Software Institute

²Microsoft Research

³Technical University of Madrid (UPM)

Many micro-architectural attacks rely on the attacker to efficiently find small *eviction sets*: addresses that map to the same cache set. We present a decisive execution primitive for cache side-channel attacks that finds small eviction sets in a systematic literature.

We begin with a probabilistic study. We begin by estimating the probability that a given address maps to a given cache set. We then use this to estimate the probability that a given address maps to a given cache set.

Accessing a large enough set of virtual addresses is sufficient for evicting any content from the cache. Such large eviction sets increase the time required for probing, and they introduce noise due to the memory accesses. For targeted and stealthy evictions, content one hence seeks to identify evictions with a fine-grained monitoring of memory accesses. We present a technique for identifying evictions with high accuracy, which is fundamental, for example, for enforcing that memory accesses are not rowhammer attacks [6]. Our technique increases the number of accesses executed by the attacker, but it also reduces the minimal set of addresses required to evict a given cache set.



Theory and Practice of Finding Eviction Sets [S&P19]



LLC Prime+Probe for VM Detection




pick eviction set(s)



preload stage + prime

VMM executes



cpuid



compute max access time among
lines in the set



Results

i7-8665U (8MB, 16 ways, 128 colors)

- VirtualBox 6.1, Windows host: 20/128
- VirtualBox 6.1, Linux host: 18/128
- VMware Workstation Pro, Windows host: 10/90
 - Eviction sets not found for every color
- QEMU-KVM 4.2.50: 13/128



Limitations

- Execution time may be long for big caches ($>$ sandboxes timeout)
- Eviction set construction may fail (e.g. non-inclusive LLC)



Questions?